

가독성 비교: 60ch vs 66ch vs 72ch

동일 콘텐츠·동일 폰트·동일 줄간격(1.84)에서 본문 최대 너비(ch 단위)만 변경하여 비교합니다.

기준: 현재 블로그 CSS `.post-body p { max-width: 66ch; line-height: 1.84; }` | WCAG 1.4.8(AAA) $\leq 80\text{ch}$ | Baymard 권장 50–75ch | Ruder 최적 50–60ch

60자

60ch

Ruder 최적 구간

Baymard 하한

Daily AI Brief: Agentic Development and Policy Distillation Lead Trends

Today's AI landscape is marked by significant activity in agentic frameworks on GitHub, alongside new research in policy distillation and community discussions on AI-driven SaaS registration.

The AI ecosystem today highlights a strong focus on agentic frameworks, with NousResearch's 'hermes-agent' leading GitHub velocity. Concurrently, new academic work on 'Trust Region Policy Distillation' is drawing attention in research circles. Community discussions are also gravitating towards practical applications, such as AI-powered SaaS registration.

Repository Momentum

Fresh GitHub projects worth scanning before the feed turns over.

- **NousResearch/hermes-agent** — The agent that grows with you. 214.2k stars, +800/7d. Updated 1h ago.
- **BerriAI/litellm** — Python SDK, Proxy Server (AI Gateway) to call 100+ LLM APIs in OpenAI format, with cost tracking, guardrails, loadbalancing and logging. 53.5k stars.
- **openai/codex** — Lightweight coding agent that runs in your terminal. 97.7k stars, +800/7d.

- **sickn33/agentic-awesome-skills** — Installable GitHub library of 1,900+ agentic skills for Claude Code, Cursor, Codex CLI, and more. 43.1k stars.
- **headroomlabs-ai/headroom** — Compress tool outputs, logs, files, and RAG chunks before they reach the LLM. 60-95% fewer tokens, same answers. 58.9k stars.

Fresh Papers

New research worth bookmarking for a deeper read.

- **Trust Region Policy Distillation** (HF Papers) — Policy distillation with trust region constraints for stable transfer.
- **Agentic Patterns Taxonomy** — Comprehensive taxonomy of agentic patterns and evaluation benchmarks.
- **Preference Distillation** — Distilling preference data into smaller models without reward modeling.

Community Wire

Discussions and show-and-tells from the community.

- **Show GN: MCP server that registers the created SaaS by simply telling AI** — GeekNews community wire.
- **Natural language to deployed service in under 2 minutes** — Demo of AI-driven service provisioning.

Technical Deep-dive

The `hermes-agent` architecture uses a reducer/effect pattern where state transitions are pure and side effects are pushed to an effect runner. This separates concerns cleanly:

```
// Reducer (pure)
function reduce(state, action) {
  switch (action.type) {
    case 'QUERY': return { ...state, query: action.payload };
    case 'RESULTS': return { ...state, results: action.payload };
  }
}

// Effect runner (impure)
async function runEffect(effect, dispatch) {
  if (effect.type === 'SEARCH') {
    const results = await searchAPI(effect.payload);
    dispatch({ type: 'RESULTS', payload: results });
  }
}
```

This pattern keeps the core logic testable and the I/O boundaries explicit. The same pattern applies to the news desk workflow where candidate selection, drafting, and publishing are separate effect types.

"The best agent frameworks don't hide complexity—they make the boundaries visible so you can reason about them."

Metrics Summary

Today's brief processed 10 repository signals, 10 paper picks, and 10 community items. The lead story combines all three modalities into a single coherent narrative.

66자 (현재)

66ch

Baymard 권장 내

WCAG 1.4.8 AAA 준수

Daily AI Brief: Agentic Development and Policy Distillation Lead Trends

Today's AI landscape is marked by significant activity in agentic frameworks on GitHub, alongside new research in policy distillation and community discussions on AI-driven SaaS registration.

The AI ecosystem today highlights a strong focus on agentic frameworks, with NousResearch's 'hermes-agent' leading GitHub velocity. Concurrently, new academic work on 'Trust Region Policy Distillation' is drawing attention in research circles. Community discussions are also gravitating towards practical applications, such as AI-powered SaaS registration.

Repository Momentum

Fresh GitHub projects worth scanning before the feed turns over.

- **NousResearch/hermes-agent** — The agent that grows with you. 214.2k stars, +800/7d. Updated 1h ago.
- **BerriAI/litellm** — Python SDK, Proxy Server (AI Gateway) to call 100+ LLM APIs in OpenAI format, with cost tracking, guardrails, loadbalancing and logging. 53.5k stars.
- **openai/codex** — Lightweight coding agent that runs in your terminal. 97.7k stars, +800/7d.

- **sickn33/agentic-awesome-skills** — Installable GitHub library of 1,900+ agentic skills for Claude Code, Cursor, Codex CLI, and more. 43.1k stars.
- **headroomlabs-ai/headroom** — Compress tool outputs, logs, files, and RAG chunks before they reach the LLM. 60-95% fewer tokens, same answers. 58.9k stars.

Fresh Papers

New research worth bookmarking for a deeper read.

- **Trust Region Policy Distillation** (HF Papers) — Policy distillation with trust region constraints for stable transfer.
- **Agentic Patterns Taxonomy** — Comprehensive taxonomy of agentic patterns and evaluation benchmarks.
- **Preference Distillation** — Distilling preference data into smaller models without reward modeling.

Community Wire

Discussions and show-and-tells from the community.

- **Show GN: MCP server that registers the created SaaS by simply telling AI** — GeekNews community wire.
- **Natural language to deployed service in under 2 minutes** — Demo of AI-driven service provisioning.

Technical Deep-dive

The `hermes-agent` architecture uses a reducer/effect pattern where state transitions are pure and side effects are pushed to an effect runner. This separates concerns cleanly:

```
// Reducer (pure)
function reduce(state, action) {
  switch (action.type) {
    case 'QUERY': return { ...state, query: action.payload };
    case 'RESULTS': return { ...state, results: action.payload };
  }
}

// Effect runner (impure)
async function runEffect(effect, dispatch) {
  if (effect.type === 'SEARCH') {
    const results = await searchAPI(effect.payload);
    dispatch({ type: 'RESULTS', payload: results });
  }
}
```

This pattern keeps the core logic testable and the I/O boundaries explicit. The same pattern applies to the news desk workflow where candidate selection, drafting, and publishing are separate effect types.

"The best agent frameworks don't hide complexity—they make the boundaries visible so you can reason about them."

Metrics Summary

Today's brief processed 10 repository signals, 10 paper picks, and 10 community items. The lead story combines all three modalities into a single coherent narrative.

Daily AI Brief: Agentic Development and Policy Distillation Lead Trends

Today's AI landscape is marked by significant activity in agentic frameworks on GitHub, alongside new research in policy distillation and community discussions on AI-driven SaaS registration.

The AI ecosystem today highlights a strong focus on agentic frameworks, with NousResearch's 'hermes-agent' leading GitHub velocity. Concurrently, new academic work on 'Trust Region Policy Distillation' is drawing attention in research circles. Community discussions are also gravitating towards practical applications, such as AI-powered SaaS registration.

Repository Momentum

Fresh GitHub projects worth scanning before the feed turns over.

- **NousResearch/hermes-agent** — The agent that grows with you. 214.2k stars, +800/7d. Updated 1h ago.
- **BerriAI/litellm** — Python SDK, Proxy Server (AI Gateway) to call 100+ LLM APIs in OpenAI format, with cost tracking, guardrails, loadbalancing and logging. 53.5k stars.
- **openai/codex** — Lightweight coding agent that runs in your terminal. 97.7k stars, +800/7d.

- **sickn33/agentic-awesome-skills** — Installable GitHub library of 1,900+ agentic skills for Claude Code, Cursor, Codex CLI, and more. 43.1k stars.
- **headroomlabs-ai/headroom** — Compress tool outputs, logs, files, and RAG chunks before they reach the LLM. 60-95% fewer tokens, same answers. 58.9k stars.

Fresh Papers

New research worth bookmarking for a deeper read.

- **Trust Region Policy Distillation** (HF Papers) — Policy distillation with trust region constraints for stable transfer.
- **Agentic Patterns Taxonomy** — Comprehensive taxonomy of agentic patterns and evaluation benchmarks.
- **Preference Distillation** — Distilling preference data into smaller models without reward modeling.

Community Wire

Discussions and show-and-tells from the community.

- **Show GN: MCP server that registers the created SaaS by simply telling AI** — GeekNews community wire.
- **Natural language to deployed service in under 2 minutes** — Demo of AI-driven service provisioning.

Technical Deep-dive

The `hermes-agent` architecture uses a reducer/effect pattern where state transitions are pure and side effects are pushed to an effect runner. This separates concerns cleanly:

```
// Reducer (pure)
function reduce(state, action) {
  switch (action.type) {
    case 'QUERY': return { ...state, query: action.payload };
    case 'RESULTS': return { ...state, results: action.payload };
  }
}

// Effect runner (impure)
async function runEffect(effect, dispatch) {
  if (effect.type === 'SEARCH') {
    const results = await searchAPI(effect.payload);
    dispatch({ type: 'RESULTS', payload: results });
  }
}
```

This pattern keeps the core logic testable and the I/O boundaries explicit. The same pattern applies to the news desk workflow where candidate selection, drafting, and publishing are separate effect types.

"The best agent frameworks don't hide complexity—they make the boundaries visible so you can reason about them."

Metrics Summary

Today's brief processed 10 repository signals, 10 paper picks, and 10 community items. The lead story combines all three modalities into a single coherent narrative.

피드백 기준: 줄 길이(문자 수) · 눈 이동 거리 · 단락 밀도감 · 모바일($\leq 560\text{px}$) 강제 1.8 줄간격과의 조화

스크린샷 저장: Cmd/Ctrl + P → PDF 저장 또는 브라우저 스크린샷 도구 사용